

ResQ-Guard

AI Video Generation Prompt Guide

A scene-by-scene prompt pack for generating a cinematic demo video (Sora / Runway / Pika / Kling / Luma) showing how ResQ-Guard solves real-world urban traffic and emergency-response problems.

Includes: master style guide, 8-scene shot list, transition bridges for seamless continuity, audio direction, text overlay suggestions, and practical generation tips.

Master Style Guide

Paste this before every individual scene prompt, or keep it saved separately to copy from each time you generate a new clip.

"Cinematic aerial and street-level footage, cool blue-and-cyan color grade, night-and-dusk urban tech aesthetic, subtle lens flares from streetlights, futuristic HUD-style overlays in cyan/white, slight film grain, smooth drone camera movement, 24fps cinematic motion blur, high contrast shadows, Indian metropolitan city setting (wide roads, flyovers, traffic signals, auto-rickshaws mixed with cars), realistic but slightly stylized like a tech product launch video, no visible text or logos on real brands."

Why This Matters

Most AI video generators cap clips at 5–20 seconds and have no memory between generations. Repeating this exact style phrase across every scene prompt is what keeps color grade, mood, and visual language consistent when you stitch separately generated clips together in an editor.

Scene-by-Scene Prompts

Scene 1 — The Problem (~10s)

Aerial drone shot descending over a dense, gridlocked city intersection at dusk. Hundreds of cars, buses, and auto-rickshaws are bumper to bumper. Horns blare, heat haze shimmers off the asphalt. Camera slowly pushes in toward a single ambulance stuck in the middle of the jam, its lights flashing uselessly. Mood: frustration, chaos, urgency. Ends on the ambulance's flashing light filling the frame, pulsing rhythmically.

Scene 2 — Surveillance Network Activates (~8s)

Extreme close-up opens matching the previous pulse's color and rhythm, then resolves into a traffic CCTV camera slowly panning across the intersection. A soft cyan scanning light sweeps across the lens. Cut to a first-person camera-vision shot of the street below, with a faint cyan grid overlay appearing on the road and vehicles, like a system waking up. Ends with the grid settled.

Transition in: cross-dissolve (~0.5s) on the pulsing light — the ambulance light 'becomes' the camera's scanning light.

Scene 3 — AI Detection & Plate Recognition (~10s)

Same static high-angle CCTV framing continues, grid overlay already active, camera now tracks right along the road, following the stuck ambulance. A thin cyan bounding box locks onto the ambulance, then a smaller box locks onto its license plate, with fine scanning lines sweeping over the plate character by character. Ends on the box confirming green/matched, with a glowing pulse radiating outward like a data ping.

Transition in: whip-pan continuation or straight cut — camera-angle continuity carries it, no dissolve needed.

Scene 4 — Command Center (~10s)

Interior shot of a dark, modern control room. A monitor screen flickers to life exactly as a green ping/notification appears on it, echoing the same glow color and pulse timing as the previous scene. Large wall-mounted screens glow with live city maps, camera feeds, and blinking alert icons in cyan and amber. A silhouetted operator taps a control surface; a green route line begins animating outward from the ambulance's icon toward a hospital icon.

Transition in: match-on-action cut — the Scene 3 'ping' visually continues as the monitor alert here.

Scene 5 — Green Corridor / Signal Preemption (~12s)

Hard cut on the operator's tap: at street level, the first traffic light in the sequence switches to green in the same instant, as if triggered by that action. Down the road, each light flips green in a wave, moving toward camera. Cut to an aerial shot of the same road, showing a glowing green corridor snaking ahead of the ambulance as it moves freely, other vehicles pulling aside. Camera rises until the city grid becomes visible from above, lights forming a network pattern.

Transition in: hard cut, no dissolve — feels causal and instant, not like a scene break.

Scene 6 — Digital Twin / Predictive Simulation (~8s)

The aerial city grid from the previous shot dissolves into the same pattern rendered as a glowing holographic wireframe hovering above a control-room table, tiny animated vehicle dots flowing along roads. A hand gestures near the hologram; one road segment highlights and reroutes, showing predicted congestion easing.

Transition in: cross-dissolve, longer (~1–1.5s) — this is the intentional 'reality → data visualization' conceptual bridge; let it breathe.

Scene 7 — Citizen & Everyday Impact (~10s)

The pulse shape and color from the hologram cut to a matching notification icon on a phone screen in a woman's hand — she is reporting a stalled vehicle via a simple app interface (clean UI, cyan accents, no real logos). Cut to the same control-room screen lighting up with her report as a new map pin. Cut to a wide shot of the ambulance from Scene 1, now arriving smoothly at a hospital entrance, staff already waiting. Camera pulls back and up into an aerial angle.

Transition in: match cut on pulse shape/color, straight cut, near-instant.

Scene 8 — Closing / Brand Moment (~8s)

Aerial pull-back continues rising, matching the previous rise speed, the hospital shrinking to a small point of light amid organized, free-flowing traffic streams like a nervous system's synapses. A soft cyan grid briefly fades in over the skyline, then fades to black.

Transition in: continuous aerial movement, matched camera speed — no cut needed if rise rate aligns; otherwise a quick cross-dissolve.

Continuity & Transition Tips

- Keep cross-dissolves short (8–15 frames at 24fps, ~0.3–0.6s) except the Scene 5→6 conceptual bridge, which can run ~1–1.5s.
- Add a subtle speed ramp (5–10% slow-mo in, ease back to normal) right at each cut point to smooth motion mismatches between separately generated clips.
- Apply one consistent color grade / LUT across all clips in your editor at the end — AI generators drift slightly clip to clip even with the same style prompt.
- Add a constant low-frequency ambient hum or drone under the entire video — audio continuity does a lot of the psychological work of making cuts feel seamless.
- If your tool supports image-to-video / frame continuation (Runway Gen-3, Kling, Luma), feed the last frame of one clip as the starting image for the next scene instead of a fresh text prompt — this is the single most effective way to guarantee visual continuity.

Audio / Music Direction

Scenes	Direction
1–2	Tense, low ambient hum, distant traffic noise, no music yet
3–4	Minimal, pulsing synth undertone builds in
5–6	Music swells, hopeful synth-driven build
7–8	Warm strings/pads layered over synth, resolving to a calm outro

Text Overlays (Add in Post)

Most AI video tools render on-screen text poorly — add these as overlays in your editor instead.

Scene	Overlay Text
Scene 1	Every minute lost in traffic can cost a life.
Scene 3	ResQ-Guard detects it in real time.
Scene 5	Clears the path — instantly.
Scene 6	Predicts congestion before it happens.
Scene 8	ResQ-Guard. Built for the city that never stops.

Practical Generation Tips

- Generate each scene as a separate clip — most tools cannot hold an 8-scene narrative in a single prompt.

- Keep the master style guide phrase at the start of every scene prompt to preserve consistent color grade and mood.
- Use 16:9 aspect ratio for hackathon/YouTube submission, or 9:16 for Instagram/LinkedIn reels.
- Regenerate a scene 2–3 times if the model hallucinates readable text or logos — none of these prompts request legible signage, but generators sometimes add gibberish on their own.